

HPA – Horizontal Pod Autoscaler

Passo a Passo

Para configurar o Horizontal Auto Scaler é preciso, primeiro, fazer o deploy do servidor que faz a coleta das métricas dos Nodes e dos Pods. Para isso:

- 1) No terminal onde está o seu cluster do Kubernetes, digite o comando “kubectl apply -f <https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml>”

```
% kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
```

- 2) Digite o comando “kubectl get pods -A” e verifique se o POD “metrics-server” está em “running”:

```
marceloortiz@MacBook-Air-de-Marcelo-2 hack % kubectl get pods -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-6d4b75cb6d-blprk	1/1	Running	2 (12m ago)	2d5h
kube-system	coredns-6d4b75cb6d-tnrtg	1/1	Running	2 (12m ago)	2d5h
kube-system	etcd-ortiz-control-plane	1/1	Running	2 (12m ago)	2d5h
kube-system	kindnet-zxpm5	1/1	Running	2 (12m ago)	2d5h
kube-system	kube-apiserver-ortiz-control-plane	1/1	Running	2 (12m ago)	2d5h
kube-system	kube-controller-manager-ortiz-control-plane	1/1	Running	2 (12m ago)	2d5h
kube-system	kube-proxy-nlm69	1/1	Running	2 (12m ago)	2d5h
kube-system	kube-scheduler-ortiz-control-plane	1/1	Running	2 (12m ago)	2d5h
kube-system	metrics-server-678f4bf65b-4wg8p	0/1	Running	0	30s
local-path-storage	local-path-provisioner-9cd9bd544-wxr54	1/1	Running	2 (12m ago)	2d5h

- 3) Verifique os eventos do seu pod “metrics-server”, caso exista alguma falha com “HTTP probe failed with statuscode: 500” siga o passo 4, caso contrário siga para o passo 6:

Para verificar os eventos do pod, digite: kubectl describe pods -n kube-system “nome do pod de métricas”:

```
kubectl describe pods -n kube-system metrics-server-678f4bf65b-4wg8p
```

O evento com erro é:

```
node.kubernetes.io/unreachable:NoExecute op=Exists for 30s
```

Events:	Type	Reason	Age	From	Message
Normal	Scheduled	5m35s	default-scheduler	Successfully assigned kube-system/metrics-server-678f4bf65b-4wg8p to ortiz-control-plane	
Normal	Pulled	5m34s	kubelet	Container image “k8s.gcr.io/metrics-server/metrics-server:v0.6.1” already present on machine	
Normal	Created	5m34s	kubelet	Created container metrics-server	
Normal	Started	5m34s	kubelet	Started container metrics-server	
Warning	Unhealthy	25s (x32 over 5m5s)	kubelet	Readiness probe failed: HTTP probe failed with statuscode: 500	

- 4) Para corrigir esse erro, será preciso editar o POD “metrics-server” e adicionar as linhas abaixo na especificação do containers:

command:

- /metrics-server
- kubelet-insecure-tls
- kubelet-preferred-address-types=InternalIP

Para isso digite “kubectl edit deploy -n kube-system metrics-server”:

```
kubectl edit deploy -n kube-system metrics-server
```

Pressione a tecla “i” e insira as linhas acima logo abaixo da linha “--metric-resolution=15s”

```
spec:
  containers:
    - args:
      - --cert-dir=/tmp
      - --secure-port=4443
      - --kubelet-preferred-address-types=InternalIP,ExternalIP,Hostname
      - --kubelet-use-node-status-port
      - --metric-resolution=15s
      command:
      - /metrics-server
      - --kubelet-insecure-tls
      - --kubelet-preferred-address-types=InternalIP
```

Cuidado com a indentação, pois as novas linhas precisam iniciar na mesma coluna das demais linhas.

Saia salvando, pressionando a tecla “esc” e depois digitando “:wq”.

- 5) Verifique se agora o POD do “metrics-server” está em execução sem nenhum erro nos últimos eventos, para isso digite “kubectl describe pods -n kube-system “nome do pod do metric server”:

```
Events:
  Type    Reason      Age   From              Message
  ----    -
  Normal  Scheduled   103s  default-scheduler Successfully assigned kube-system/metrics-server-69cb89695d-t5ll5 to ortiz-control-plane
  Normal  Pulled      103s  kubelet           Container image "k8s.gcr.io/metrics-server/metrics-server:v0.6.1" already present on machine
  Normal  Created     103s  kubelet           Created container metrics-server
  Normal  Started     103s  kubelet           Started container metrics-server
```

- 6) Veja que agora é possível executar os comandos “kubectl top pods -A” e “kubectl top nodes”. Esse comando vai trazer a utilização de CPU e memória dos seus PODs e Nodes.

```
marceloortiz@MacBook-Air-de-Marcelo-2 hack % kubectl top nodes
NAME                                CPU(cores)   CPU%   MEMORY(bytes)   MEMORY%
ortiz-control-plane                 240m         6%     747Mi           18%
```

```
marceloortiz@MacBook-Air-de-Marcelo-2 hack % kubectl top pods -A
NAMESPACE      NAME                                CPU(cores)   MEMORY(bytes)
kube-system    coredns-6d4b75cb6d-blprk           7m           12Mi
kube-system    coredns-6d4b75cb6d-tnrtg           6m           12Mi
kube-system    etcd-ortiz-control-plane            30m          48Mi
kube-system    kindnet-zxpm5                       2m           10Mi
kube-system    kube-apiserver-ortiz-control-plane   82m          319Mi
kube-system    kube-controller-manager-ortiz-control-plane 31m          42Mi
kube-system    kube-proxy-nlm69                    1m           10Mi
kube-system    kube-scheduler-ortiz-control-plane   6m           16Mi
kube-system    metrics-server-69cb89695d-t5ll5     10m          18Mi
local-path-storage local-path-provisioner-9cd9bd544-wxr54 3m           6Mi
```

Agora já podemos fazer o deploy de uma aplicação qualquer. No nosso caso vamos fazer um deploy de um nginx. Essa aplicação que vamos fazer a configuração do HPA.

- 7) Crie um arquivo com o nome “deploy.yaml” com as informações abaixo:

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: scaling-deploy
spec:
  replicas: 1
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx
        resources:
          requests:
            memory: "50Mi"
            cpu: "500m"

```

- 8) Faça o deploy desse arquivo com o comando “kubectl apply -f deploy.yaml”:

```

marceloortiz@MacBook-Air-de-Marcelo-2 k8s % kubectl apply -f deploy.yaml
deployment.apps/scaling-deploy created

```

- 9) Verifique o consumo desse POD com o comando “kubectl top pods”:

```

marceloortiz@MacBook-Air-de-Marcelo-2 k8s % kubectl top pods
NAME                                CPU(cores)   MEMORY(bytes)
scaling-deploy-8f458dc5b-4z699      0m           6Mi

```

Veja que o consumo de CPU está bem baixo.

Agora vamos criar nosso HPA com base na utilização da CPU, para isso:

- 10) Crie um arquivo chamado “hpa.yaml” com o conteúdo abaixo:

```

apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: scaling-deploy
spec:
  maxReplicas: 10
  minReplicas: 1
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: scaling-deploy
  targetCPUUtilizationPercentage: 70

```

- 11) Faça o deploy do HPA com o comando “kubectl apply -f hpa.yaml”:

```
marceloortiz@MacBook-Air-de-Marcelo-2 k8s % kubectl apply -f hpa.yaml
horizontalpodautoscaler.autoscaling/scaling-deploy created
```

- 12) Verifique se o HPA foi criado corretamente com o comando “kubectl get hpa”:

```
marceloortiz@MacBook-Air-de-Marcelo-2 k8s % kubectl get hpa
NAME          REFERENCE                TARGETS  MINPODS  MAXPODS  REPLICAS  AGE
scaling-deploy  Deployment/scaling-deploy  1%/70%   1         10        1          2m55s
```

Agora vamos simular a carga no nosso POD e verificar o comportamento do HPA. Para isso:

- 13) Conecte no seu POD com o comando “kubectl exec -it “nome do pod” -- bash”:

```
% kubectl exec -it scaling-deploy-6686b77c59-9mj9s -- bash
```

- 14) Dentro do seu POD, execute o comando “apt update”:

```
root@scaling-deploy-6686b77c59-9mj9s:/# apt update
Get:1 http://deb.debian.org/debian bullseye InRelease [116 kB]
Get:2 http://deb.debian.org/debian-security bullseye-security InRelease [44.1 kB]
Get:3 http://deb.debian.org/debian bullseye-updates InRelease [44.1 kB]
Get:4 http://deb.debian.org/debian bullseye/main arm64 Packages [8069 kB]
Get:5 http://deb.debian.org/debian-security bullseye-security/main arm64 Packages [165 kB]
Get:6 http://deb.debian.org/debian bullseye-updates/main arm64 Packages [2600 B]
Fetched 8441 kB in 4s (2045 kB/s)
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
All packages are up to date.
```

- 15) Dentro do seu POD, execute o comando “apt install stress -y”:

```
root@scaling-deploy-6686b77c59-9mj9s:/# apt install stress -y
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following NEW packages will be installed:
  stress
0 upgraded, 1 newly installed, 0 to remove and 0 not upgraded.
Need to get 21.6 kB of archives.
After this operation, 54.3 kB of additional disk space will be used.
Get:1 http://deb.debian.org/debian bullseye/main arm64 stress arm64 1.0.4-7 [21.6 kB]
Fetched 21.6 kB in 3s (7725 B/s)
debconf: delaying package configuration, since apt-utils is not installed
Selecting previously unselected package stress.
(Reading database ... 7815 files and directories currently installed.)
Preparing to unpack .../stress_1.0.4-7_arm64.deb ...
Unpacking stress (1.0.4-7) ...
Setting up stress (1.0.4-7) ...
```

- 16) Dentro do seu POD, execute o comando “stress --cpu 1”:

```
root@scaling-deploy-6686b77c59-9mj9s:/# stress --cpu 1
stress: info: [318] dispatching hogs: 1 cpu, 0 io, 0 vm, 0 hdd
```

- 17) Abra um outro terminal no seu cluster Kubernetes e verifique o consumo do POD com o comando “kubectl top pods”:

```
MacBook-Air-de-Marcelo-2:k8s marceloortiz$ kubectl top pod
NAME                                CPU(cores)   MEMORY(bytes)
scaling-deploy-6686b77c59-6941s    6m           3Mi
scaling-deploy-6686b77c59-9mj9s    999m         32Mi
```

Veja que o POD está consumindo 999m de CPU e o HPA começou a criar novos PODs.

- 18) Verifique quantos nos PODs foram criados automaticamente com o comando “kubectl get pods”:

```
MacBook-Air-de-Marcelo-2:k8s marceloortiz$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
scaling-deploy-6686b77c59-6941s    1/1     Running   0           3m29s
scaling-deploy-6686b77c59-9mj9s    1/1     Running   0           10m
scaling-deploy-6686b77c59-pjxkb    1/1     Running   0           3m29s
```

No caso dos nossos testes foram criados 2 novos PODs.

- 19) Interrompa o “stress” que está sendo executado no primeiro pod digitando as teclas “ctrl + c”. Aguarde até que o HPA verifique que o consumo voltou abaixo dos 70% de CPU e observe se os PODs que foram criados durante o pico de consumo foram destruídos.

- 20) Para limpar o ambiente:

- Remover o HPA: `kubectl delete -f hpa.yaml`
- Remover o nginx: `kubectl delete -f deploy.yaml`
- Remover o servidor de métricas: `kubectl delete -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml`